

penguin Classifier

```
library(tidyverse)
theme_set(theme_bw())
```

Data

```
library(palmerpenguins)
penguins
```

```
# A tibble: 344 x 8
  species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>   <fct>         <dbl>         <dbl>           <int>         <int>
1 Adelie  Torgersen         39.1          18.7            181          3750
2 Adelie  Torgersen         39.5          17.4            186          3800
3 Adelie  Torgersen         40.3           18             195          3250
4 Adelie  Torgersen          NA           NA              NA           NA
5 Adelie  Torgersen         36.7          19.3            193          3450
6 Adelie  Torgersen         39.3          20.6            190          3650
7 Adelie  Torgersen         38.9          17.8            181          3625
8 Adelie  Torgersen         39.2          19.6            195          4675
9 Adelie  Torgersen         34.1          18.1            193          3475
10 Adelie Torgersen         42           20.2            190          4250
# i 334 more rows
# i 2 more variables: sex <fct>, year <int>
```

Explore Data

```
penguins |> count(species)
```

```
# A tibble: 3 x 2
  species      n
  <fct>    <int>
1 Adelie  222
2 Adelie  134
```

```

  <fct>      <int>
1 Adelie      152
2 Chinstrap   68
3 Gentoo     124

```

```
penguins |> count(island)
```

```

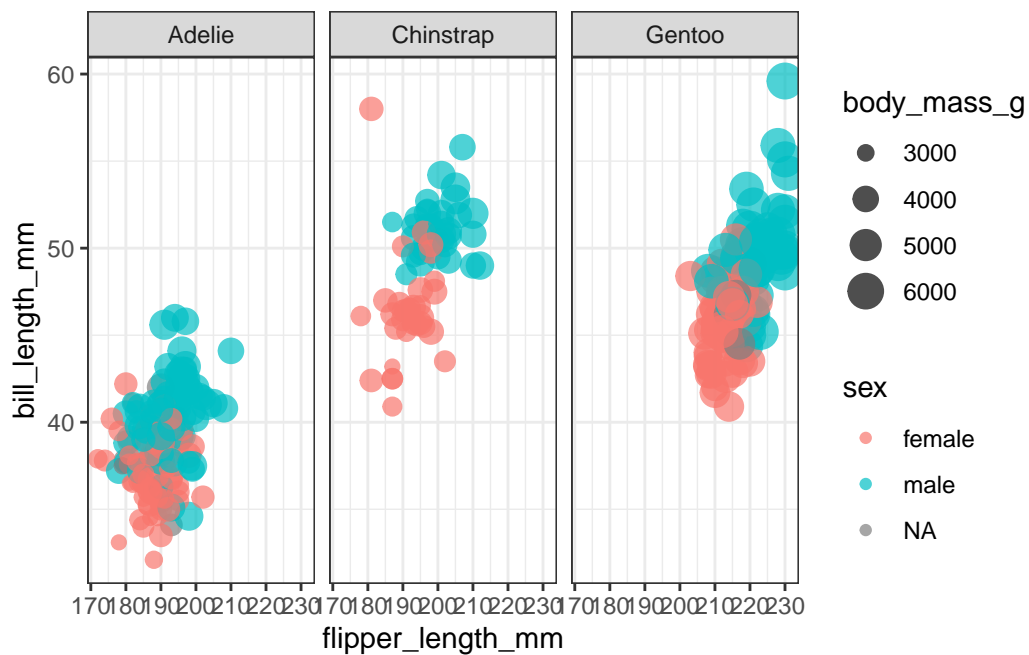
# A tibble: 3 x 2
  island      n
  <fct>      <int>
1 Biscoe     168
2 Dream      124
3 Torgersen   52

```

```

penguins |>
  ggplot(aes(flipper_length_mm, bill_length_mm, color=sex, size=body_mass_g))+
  geom_point(alpha=0.7)+
  facet_wrap(~species)

```



```

penguins_df <- penguins %>%
  filter(!is.na(sex)) %>%
  select(-year, -island)
penguins_df

```

```
# A tibble: 333 x 6
```

	species	bill_length_mm	bill_depth_mm	flipper_length_mm	body_mass_g	sex
	<fct>	<dbl>	<dbl>	<int>	<int>	<fct>
1	Adelie	39.1	18.7	181	3750	male
2	Adelie	39.5	17.4	186	3800	female
3	Adelie	40.3	18	195	3250	female
4	Adelie	36.7	19.3	193	3450	female
5	Adelie	39.3	20.6	190	3650	male
6	Adelie	38.9	17.8	181	3625	female
7	Adelie	39.2	19.6	195	4675	male
8	Adelie	41.1	17.6	182	3200	female
9	Adelie	38.6	21.2	191	3800	male
10	Adelie	34.6	21.1	198	4400	male

i 323 more rows

Build a model

```
library(tidymodels)

set.seed(123)
penguin_split <- initial_split(penguins_df, strata=sex)
penguin_train <- training(penguin_split)
penguin_test <- testing(penguin_split)

set.seed(234)
penguin_boot <- bootstraps(penguin_train)
penguin_boot
```

```
# Bootstrap sampling
# A tibble: 25 x 2
  splits          id
  <list>         <chr>
1 <split [249/92]> Bootstrap01
2 <split [249/99]> Bootstrap02
3 <split [249/94]> Bootstrap03
4 <split [249/97]> Bootstrap04
5 <split [249/93]> Bootstrap05
6 <split [249/90]> Bootstrap06
7 <split [249/86]> Bootstrap07
8 <split [249/92]> Bootstrap08
9 <split [249/90]> Bootstrap09
10 <split [249/89]> Bootstrap10
# i 15 more rows
```

```
glm_spec <- logistic_reg()%>%
  set_engine("glm")

rf_spec <- rand_forest() %>%
  set_mode("classification") %>%
  set_engine("ranger")
```

```
penguin_wf <- workflow() %>%
  add_formula(sex ~ .)

penguin_wf
```

```
== Workflow =====
Preprocessor: Formula
Model: None

-- Preprocessor -----
sex ~ .
```

```
glm_rs <- penguin_wf %>%
  add_model(glm_spec) %>%
  fit_resamples(
    resamples = penguin_boot,
    control=control_resamples(save_pred=TRUE, verbose=TRUE)
  )

rf_rs <- penguin_wf %>%
  add_model(rf_spec) %>%
  fit_resamples(
    resamples=penguin_boot,
    control=control_resamples(save_pred=TRUE, verbose=TRUE)
  )
```

Evaluate modeling

```
collect_metrics(rf_rs)
```

```
# A tibble: 3 x 6
  .metric      .estimator  mean    n std_err .config
  <chr>        <chr>      <dbl> <int>  <dbl> <chr>
1 accuracy    binary      0.925   25 0.00467 Preprocessor1_Model11
2 brier_class binary      0.0604  25 0.00222 Preprocessor1_Model11
3 roc_auc     binary      0.980   25 0.00188 Preprocessor1_Model11
```

```
collect_metrics(glm_rs)
```

```
# A tibble: 3 x 6
```

	.metric	.estimator	mean	n	std_err	.config
	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	accuracy	binary	0.917	25	0.00481	Preprocessor1_Model11
2	brier_class	binary	0.0603	25	0.00388	Preprocessor1_Model11
3	roc_auc	binary	0.977	25	0.00471	Preprocessor1_Model11

```
glm_rs %>%
  conf_mat_resampled()
```

```
# A tibble: 4 x 3
```

	Prediction	Truth	Freq
	<fct>	<fct>	<dbl>
1	female	female	40.8
2	female	male	3.8
3	male	female	3.76
4	male	male	43.1

```
glm_rs %>%
  collect_predictions() %>%
  group_by(id) %>%
  roc_curve(sex, .pred_female)%>%
  autoplot()
```

